

# Estimating “Contact Zones” for NFL Defenders to Improve Tackling Tracking

Dennison Jackson  
Department of Statistics, Kansas State University

May 10, 2024

## Abstract

This paper introduces a method for estimating “contact zones” around defenders in American football using NFL Next Gen Stats, with the goal of opening up the field for future research into tackling. Currently, there is not much research being done on tackling using player data. Part of the issue is that it can be difficult to discern when a tackle begins in the data, especially in cases where there are multiple tacklers on a play. We focus on establishing zones around defenders such that if a ball-carrier enters a defender’s zone, it is reasonable to assume that contact has initiated. We develop several metrics with the goal of figuring out an optimal radius for our zones that aligns with real-world tackle events. Our findings suggest that a zone with a radius of 1.2 yards around defenders offers a reliable estimate for initiating contact. The implementation of these zones will help advance future research by allowing for better analysis of individual tackles as well as the ability to analyze multiple tacklers on the same play.

*Keywords:* American football, tackling, defense analysis, tracking data

# 1 Introduction

For the past six years, the National Football League (NFL) has hosted a “Big Data Bowl” contest on the data science competition platform Kaggle ([Big Data Bowl n.d.](#), [NFL Big Data Bowl 2024 n.d.](#)). Each competition has its own theme that revolves around some aspect of American football; in 2024, this theme was tackling. Contestants were asked to come up with new metrics or coaching strategies revolving around tackling by analyzing the NFL’s Next Gen Stats ([NFL Next Gen Stats n.d.](#)). This data consists of player tracking data, which includes coordinates of every player along with their orientation, speed, and acceleration among numerous other statistics which are collected at a frequency of ten frames a second for every play in every game. The NFL published Next Gen Stats for plays involving tackles for the first nine weeks for the 2022-2023 season onto Kaggle, along with supplementary datasets that include detailed information about individual plays, games, players, and tackles.

The end goal of the project is open-ended, but one of the main tracks involved the creation of metrics based around tackling. Originally, the goal of this paper was to follow this track and develop our own new statistic, but along the way, we ran into a question that shifted our perspective and led to a new project entirely: “When does a tackle start in the tracking data?” Most metrics we were considering dealt with analyzing what happens during a tackle, and in order to calculate anything related to that we need to know roughly when a tackle begins and ends.

When first glancing at the data, this may not be an obvious issue. In the tracking data, there is a variable labeled event, which consists of labels for significant events that happen during a play. Most of the values are missing, but on certain frames you may see

“pass\_caught”, “ball\_snap”, or “fumble” among others. Two of those others include “tackle” and “first\_contact.” A frame with an event of “tackle” provides a clear end to the tackle, as that is the moment the ball-carrier was officially marked as being down, meaning the play is over. “first\_contact” is a different story—it represents the first moment in a play that any defender made contact with a ball-carrier. This may seem like a good indicator of when a tackle starts, but there are some major issues with using it as such.

First contact happening does not guarantee a tackle happened from that contact. It is not unusual for a ball-carrier to be contacted by a defender and a) break free of the tackle or b) be tackled shortly after by a different defender. Both of these cases could cause serious issues with our metrics if we treat first contact as the beginning of any tackle. Even if we assumed that to be true, there is another glaring issue: what happens when there are multiple tacklers on a play? A tackle is not limited to one player, multiple tacklers can be credited as having made one on the same play. The event “first\_contact” cannot account for this, it only tells us when the first defender first made contact, but nothing about any of the other defenders.

In this paper, we will be defining a new way to identify when a tackle begins using the tracking data. We will do this by creating a “contact zone” around certain defenders, such that if the ball-carrier is within a defender’s zone, it is reasonable to assume that the defender is making contact with them. If we can estimate the size of the zone well, we will be able to get an approximation of when any defender starts making contact with the ball-carrier for any given play. This opens up the door for future research if we are wanting to study anything across the duration of a tackle or any case where we want to analyze the tackling of multiple defenders on a play.

## 1.1 Previous Work

As NFL Next Gen Stats are relatively new, especially to the public, there is not much research involving them. There’s even less research around its applications on the defensive side of the field; most papers tend to focus on the offense. [Yurko et al. \(2020\)](#) explore using Next Gen Stats to continuously model game outcomes such as the expected number of yards a ball-carrier will make it on a play. While the focus of the analysis is on the ball-carrier, they do derive measures relating to the defenders, such as how close defenders are to the ball-carrier at every point in the play. They also use Voronoi tessellations to divide the entire field into zones, where each player has a region representing the area on the field closest to them.

[Quang Nguyen & Matthews \(2024\)](#) uses player tracking data create a new metric called STRAIN that evaluates the effectiveness of a defense’s pass rush on how effectively they can break through the offensive line and put pressure on the quarterback. This analysis does not involve tackling or contact, but rather looks at the speeds of defenders and their distance to the quarterback. [Eager & Seth \(2023\)](#) also uses the Next Gen Stats to develop metrics centered around linebackers, though these do not involve tackling either.

While there has been extensive research done on tracking data from other sports such as basketball or soccer/football, there are not many sports with this data that are contact sports in the same way American football is. [Gudmundsson & Horton \(2017\)](#) has a great overview of applications of tracking data in a wide variety of sports. [Fernández et al. \(2018\)](#) uses football (soccer) tracking data to analyze how players control and influence the space around them.

Not much research has been done into defining zones for physical contact around players.

Motivated by the COVID-19 pandemic, [Gonçalves et al. \(2020\)](#) has a paper that uses football (soccer) tracking data to identify contact between players that potentially causes exposure to viruses. They monitor the amount of time players and referees are within 2m of each other with the goal of judging whether or not there is a risk of a disease spreading.

## 2 Methodology

In this paper, we choose to only estimate contact zones for all defensive backs and linebackers, from here on referred to collectively as defenders of interest (DOI). This includes all defenders except for defensive linemen. One major factor for this decision has to do with blockers. In any play, the offensive players that do not have the ball are usually trying to block the defenders to stop them from reaching the ball-carrier. While players of any position can be blocked, this tends to happen more to defensive linemen than any other defender, as right when the play starts, they immediately start grappling with the offensive linemen. The idea of a contact zone assumes that there is nothing impeding the defender from reaching the ball-carrier, so when a ball-carrier runs through a gap in the line, it is highly likely they will enter the contact zones for linemen, but the linemen will not be able to reach them to do another player being in-between them. Since my zones do not account for blockers, and defensive linemen tend to be blocked the most, we excluded them from this analysis and have left that for future work. Since linemen also have a different role to play than other defenders (they are focused on sacking the quarterback and preventing primarily runs as opposed to passes), their contact zone may differ in size from the ones assigned to a DOI, and should potentially be estimated separately.




Figure 1: Example starting frame of a play with theoretical contact zones show in gray around defenders of interest. The ball is represented by a brown circle.

## 2.1 Data Preparation

For this analysis, we decided to only consider data from the first week of the 2022-2023 NFL season, but everything done here is scalable and could be easily applied to more seasons at the expensive of a higher runtime. To prepare the data for analysis, we create four new variables:

- `gamePlayId`: a combination of the `gameId` and `playId` for a play in the form `gameId-playId`
- `doi_ind`: A dummy variable with a value of 1 if the player is a defender of interest, and 0 if not.
- `bc_ind`: A dummy variable with a value of 1 if the player is the ball-carrier, and 0 if not.

- `bcInZone`: A dummy variable with a value of 1 if the distance from the ball-carrier to a defender of interest is less than the proposed radius of the contact zone, and a 0 if not.

We then filter the plays to only include those where a tackle was performed by at least one defender of interest. Next, we extract the ball-carrier’s information from each play, namely their  $x$  and  $y$  coordinates on the field. Then, we join these coordinates back with each of the DOIs coordinates to calculate the Euclidean distance to the ball carrier at every frame on a play.

We also apply a filter to the `bcInZone` indicator if the play involves a pass. For any play with the events “`pass_arrived`” or “`pass_caught`” we set the value of `bcInZone` to 0 until the pass has arrived. We need to do this because otherwise the ball-carrier is treated as an active ball-carrier before they have even recieved the ball. This means that as they are running their route, they may be passing through DOI’s contact zones, which can throw off our results and indicate that there were more players active in the tackle than there actually were. It is illegal to start tackling a player before they have the ball, so until then no zones should be activated. Once they are in possession of the ball, they are fair game to be tackled and thus the indicator is reactivated.

After this, the preliminary data filtering is done and we can start work on evaluating the size of a contact zone.

## 2.2 Evaluating a contact zone

We define a contact zone as a circle with radius  $r$  centered around a defender. We wish to figure out the optimal value for  $r$ , such that the size of the zone accurately approximates the start of a tackle. Say, for instance, we decide on a radius of  $r = 1.5$  yards. This would

mean that whenever a ball-carrier is 1.5 yards away from a defender of interest, we want to assume they are able to begin making contact at approximately that point. In order to evaluate how good of an approximation a zone with radius  $r$  is, we propose five metrics that can potentially indicate if one radius is better than another.

Our first two metrics go hand-in-hand; we call them `meanEntryDiff` and `medianEntryDiff`. These represent the average and median differences between the frame the event “first\_contact” occurred on and the first frame that the ball-carrier entered a contact zone on a play. So for instance, if the difference between first contact and the first entry into a contact zone is 10, that means our zone is saying contact started 10 frames (or one second) before the data claims it started. A value of -5 would indicate contact started 5 frames (or half a second), and ideally we would get a difference of 0, indicating that the ball-carrier entered the zone at the same moment first contact is said to have occurred.

There is an important assumption we are making here: the event “first\_contact” is the true moment of first contact on a play. Now, is this a good assumption make? The labels were added by humans going through the plays frame by frame, and humans are prone to error. We do not believe error is a major issue here, as it should not be difficult in the majority of cases to identify where contact starts within at least few frames. A bigger issue arises when first contact is made not by a defender of interest, but by a linemen. As mentioned in the introduction, “first\_contact” does not guarantee a tackle is made. There are many plays where a linemen may rush the quarterback and contacts him (so “first\_contact” occurs), but the quarterback escapes and runs into a contact zone some time later; that is what can throw off these calculations. This means that for `mean` and `medianEntryDiff` to hold, we also want to assume that the player making first contact is a DOI.

The next three stats we propose all revolve around how well our zone identifies who was



credited as making a tackle on a play. These may be more helpful when we are talking about plays involving multiple tacklers, to see if the zones can accurately capture those occurrences. There is another assumption we need to make first: if a ball-carrier entered a defender's contact zone any point during a play, we say that defender was involved in the tackle. There are violations of this of course, a player could miss a tackle for instance. Generally speaking, there are not a large amount of missed tackles though, and even if the defender who "missed" did not get the final tackle, it is also possible they slowed the player down enough so that the actual tackle could happen, and so they still played a role. We also assume that the list of officially credited tacklers is accurate, and that all the players on that list were the only ones that were involved in the tackle. I do not love this assumption as it is not unreasonable to say that a player who missed a tackle but slowed the ball-carrier did have some role to play in the tackle. If we make this assumptions though, then we can generate the following metrics:

- `correctPct`: The proportion of plays where our predicted list of tacklers matches exactly with the official list of credited tacklers on a play.
- `noneMissingPct`: The proportion of plays where our predicted list of tacklers includes all of the officially credited tacklers (doesn't have to match exactly, could include extra players).
- `extraPct`: The proportion of plays where we predict at least one extra player is involved in the tackle that is not officially credited as a tackler.

To calculate these statistics, we count the number of zones the ball-carrier enters in each play along with a list of all the IDs of those zones' defenders. These can then be compared to the officially credited tacklers, found in a separate dataset that was included with the tracking data.

Conceptually, the first two metrics about frame differences from first contact to first zone entry seem the most useful here as they directly revolve around contact, and we are trying to estimate the size of a zone that approximates contact. The latter three metrics are more supplementary, as they do not directly involve contact. We will see that for almost any zone size, we tend to overestimate the number of tacklers involved on a play, but as stated earlier it is possible that the NFL underestimates the number of defenders important to a tackle.

## 3 Results

### 3.1 Numeric Evaluation of Zone Sizes

All of the above steps were condensed into one R function we call `evalContactZones()` that takes NFL tracking data and a proposed radius  $r$  as arguments. It returns two dataframes: `results`, which contain all five of the proposed metrics for a zone of radius  $r$ , and `defInfo`, which contains our transformed tracking data on each play, and is used primarily for visualizing plays. To test different zone sizes, we iterate over the radii  $r = 0.9, 0.95, 1.0, \dots, 2.2$ , plugging each radius into our function to be evaluated. Keep in mind that the tackles have been filtered to only include tackles where a DOI was involved, defensive linemen have been removed from the official tackles for this analysis.

Figure 2 shows the values for `meanEntryDiff` and `medianEntryDiff` for different zone sizes. The optimal radius appears to be 1.2 yards, as both the mean and median are centered at 0. This means that on average, when defenders of interest have a zone with a radius of 1.2 yards around them, the first time the ball-carrier enters a zone is the same as the event “first\_contact” is recorded, which we are assuming is the true moment of first contact. If

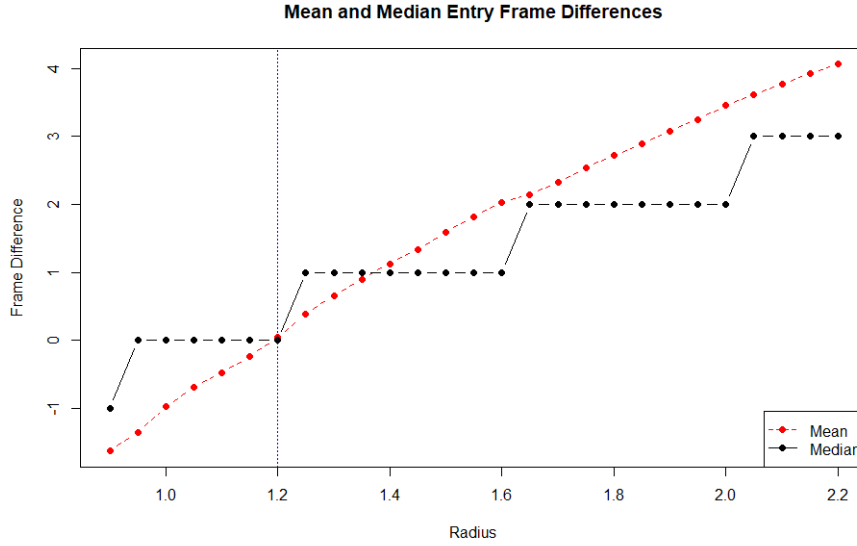


Figure 2: The mean and median frame differences for different radii. The dashed blue line at 1.2 signifies the optimal value.

1.2 yards ends up being too small of a radius, 1.35 and 1.65 yard radii seem like decent choices too as they look to have close to symmetric frame differences and are off by about 1 to 2 frames respectively on average (meaning the ball-carrier is entering the zone *before* first\_contact).

Figure 3 tells us the proportion of plays where we correctly identify every officially credited tackler. We see that 1.25 and 1.3 yards are both very similar and have the largest proportions, meaning that zones with those radii miss the least amount of official tacklers. This is very similar to the optimal value of 1.2 yards from Figure 2, and we can see that 1.2 also has one of the largest proportions here. These results along with the previous ones indicate that the optimal radius is likely around the 1.2 to 1.3 yard range.

We will not be showing plots for correctPct or extraPct here as their results were not useful and had little value for interpretation. They both exhibited highly linear trends, with correctPct decreasing and extraPct increasing as radius increases. This makes sense from

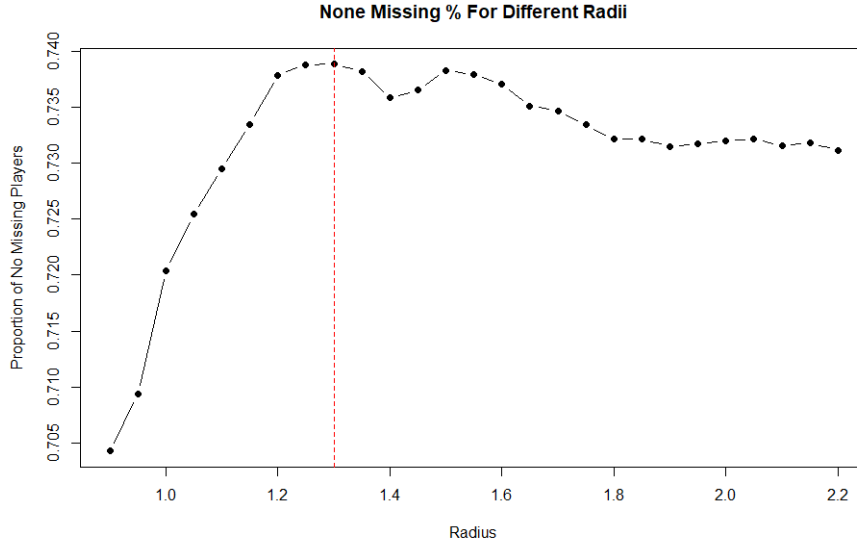


Figure 3: The proportion of plays where no officially credited tacklers are missed in our predicted tackler list. The red dashed line indicates the radius with the maximum proportion.

a logical standpoint—most plays only have one or two tacklers involved in the actual tackle, and the larger the zone size, the more zones the ball-carrier will enter in a play. Thus, correctPct will decrease as for large zones we are often adding far more players than are typically credited, and extraPct will of course increase for the same reason.

Of all the proposed metrics, we have the most faith in the mean/median frame differences as they are the stats most directly related to the actual moment of contact. There may be issues with the “first\_contact” event, but it is the closest thing we have in the data to a definitive moment of contact, so a stat that is directly involved with “first\_contact” seems more reliable for estimating a contact zone. With this in mind, we believe that 1.2 is the optimal radius for a contact zone for defensive backs and linebackers.

## 3.2 Visual Inspection of Zone Sizes

After evaluating a zone size numerically, we wrote a function called `plotGames()` that allows the user to visualize an animation of any play in our filtered data with contact zones of radius  $r$ . To evaluate our zones, we would look at the animation and then compare to replays of the real plays to see if our zones seemed to accurately reflect the play. PDFs cannot support the playback of .gif or .mp4 files, so Figure 4 shows several frames from one of the sampled games.

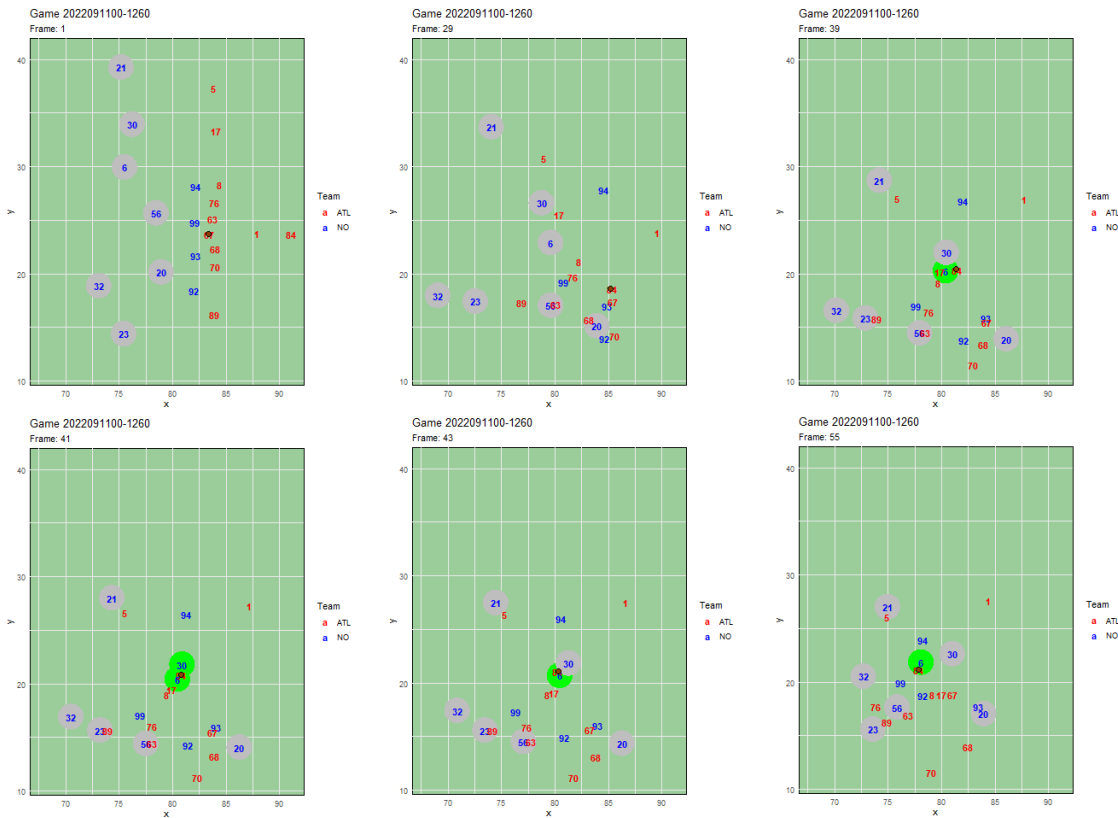


Figure 4: Six frames from a play in the ATL vs NO Week 1 game in 2022. Top-right image is first contact, where the ball-carrier (number 84) first comes into contact with a DOI (number 6). Bottom-left shows another DOI (number 30) who comes in and briefly makes contact. The play ends on frame 55 in the bottom-right with 6 finishing the tackle.

If you compare this game to the [real game footage](#) (the play in question starts at 41:52),

you can see that the zones seem to coincide well with the plays. Player 6 and player 30 both make contact at almost the same point. 30 dives at the ball-carrier and hits them before falling to the side. 6 wraps around the ball-carrier and ends up actually bringing them down. This is reflected very accurately in our data, as 30 only comes into contact for a couple frames (when he dives), but 6 stays with him all the way through first contact. If the circle was smaller, it is possible the ball-carrier would never have entered 30's zone, and if it was bigger then they may have stayed in 30's zone as even though 30 fell to the side, they still are relatively close the ball-carrier.

This is just one play of many, but across the games we watched for the most part a zone with a radius of 1.2 yards seems approximately accurate to what happens in a real play. It's not perfect of course, but it does seem to be a good estimate.

## 4 Discussion

The main goal of this project was to make future work easier for any researchers interested in analyzing tackling data. Currently, there are significant issues with identifying when a tackle begins in the tracking data, as the "first\_contact" event has issues such as not being able to account for multiple tacklers on a play. With our recommendation of a zone that is centered around each defensive back and linebacker with a radius of 1.2 yards, we provide an alternative way to measure the beginning and end of a tackle that is able to handle multiple tacklers on a play.

Future research that could benefit from incorporating these contact zones include any tackling analysis that revolves around analyzing multiple defenders within a play. For instance, one could take the idea of dividing credit among offensive players for a play from

[Yurko et al. \(2019\)](#) and apply that to defenders. Credit could be assigned through metrics such as who had the ball-carrier in their zone the longest or who caused the biggest shift in velocity of the ball-carrier. We leave the rest of the details to future researchers who may be interested.

Other research that could benefit from the zones are analyses around individual tackles, specifically about what happens across the timespan of a tackle. We could foresee potential metrics that may be able to be derived from this sort of analysis, such as which defenders can minimize expected yards after contact or create the largest shifts in momentum of the ball-carrier (i.e. who can consistently bring down a ball-carrier the fastest?).

The data the NFL provides for the competition comes with a variable called `pff_missedTackles`, that is a dummy variable that indicates if a player missed a tackle on a play. The data is created by Pro Football Focus, an organization that does a lot of work with American football statistics. The zones in this paper could provide an alternative way to assess missed tackles judging by when a ball-carrier enters and leaves a zone without being tackled. This could also be reframed as missed tackle opportunities, where we can analyze situations where the ball-carrier was in a player's zone but they did not make contact. We could see which players capitalize the most on having a ball-carrier in their zone versus those who cannot seem to make contact or finish a tackle.

## **4.1 Future Work Involving contact zones**

There is still much work to be done concerning the analysis in this paper as well. Our chief concern is the lack of contact zones for defensive linemen. One could just assign them a zone of radius 1.2 yards like the other defenders, but we think more research needs to be done to incorporate blockers before adding zones for linemen. It's also possible that linemen

need to be analyzed separately, as the zone that best fits defenders who are primarily pass defense may not work best for those who are mainly focused on stopping runs and sacking quarterbacks. In fact, it may be worth looking into separating defensive backs and linebackers and estimating separate zones, or potentially going even further and creating dynamic zones for each individual player. This would require a lot more effort, but at the moment we are assuming everybody has the same “reach” but that may not be an accurate reflection of reality.

Something else that is not an accurate reflection of reality is an assumption we have been holding throughout this entire paper: that the contact zone is a circle. This is actually not physically possible in real life, as a defender cannot tackle somebody who is behind them. Realistically, the zone should be a semi-circle or a cone directly in front of the defender, which should not be too difficult to create given that the tracking data includes the orientation of the players. While a circle is not a good approximation of a real contact zone, we do not believe it ruins this analysis or causes much error as throughout a play, it is very rare for a defender to ever have their back to an active ball-carrier, especially if they are close enough to each other that we would consider them to potentially be in contact with each other. Still, updating the zone to a more realistic shape will most likely offer some improvements to the analysis, if not just making it more reflective of the real world.

## References

*Big Data Bowl* (n.d.).

**URL:** <https://operations.nfl.com/gameday/analytics/big-data-bowl>

Eager, E. & Seth, T. (2023), ‘Investigating trade-offs made by american football linebackers



using tracking data’, *Journal of Quantitative Analysis in Sports* **19**(3), 171–185.

**URL:** <https://doi.org/10.1515/jqas-2022-0091>

Fernández, J., Bornn, L. & Fraser, S. (2018), Wide open spaces: A statistical technique for measuring space creation in professional soccer.

**URL:** <https://api.semanticscholar.org/CorpusID:209898212>

Gonçalves, B., Mendes, R., Folgado, H., Figueiredo, P., Travassos, B., Barros, H., Campos-Fernandes, A., Beckert, P. & Brito, J. (2020), ‘Can tracking data help in assessing interpersonal contact exposure in team sports during the covid-19 pandemic?’, *Sensors* **20**(21).

**URL:** <https://www.mdpi.com/1424-8220/20/21/6163>

Gudmundsson, J. & Horton, M. (2017), ‘Spatio-temporal analysis of team sports’, **50**(2).

**URL:** <https://doi.org/10.1145/3054132>

*NFL Big Data Bowl 2024* (n.d.).

**URL:** <https://www.kaggle.com/competitions/nfl-big-data-bowl-2024/data>

*NFL Next Gen Stats* (n.d.).

**URL:** <https://operations.nfl.com/gameday/technology/nfl-next-gen-stats/>

Quang Nguyen, R. Y. & Matthews, G. J. (2024), ‘Here comes the strain: Analyzing defensive pass rush in american football with player tracking data’, *The American Statistician* **78**(2), 199–208.

**URL:** <https://doi.org/10.1080/00031305.2023.2242442>

Yurko, R., Matano, F., Richardson, L. F., Granered, N., Pospisil, T., Pelechrinis, K. & Ventura, S. L. (2020), ‘Going deep: models for continuous-time within-play valuation of game outcomes in american football with tracking data’, *Journal of Quantitative Analysis*

*in Sports* **16**(2), 163–182.

**URL:** <https://doi.org/10.1515/jqas-2019-0056>

Yurko, R., Ventura, S. & Horowitz, M. (2019), ‘nflwar: a reproducible method for offensive player evaluation in football’, *Journal of Quantitative Analysis in Sports* **15**(3), 163–183.

**URL:** <https://doi.org/10.1515/jqas-2018-0010>